

```
using Pkg
Pkg.activate(".")
using BenchmarkTools
using LinearAlgebra
```

```
Activating environment at `~/Documents/schoolwork-codes/physics-215-julia/session-3/Project.toml`
```

Session 3: Types, type inferencing, and type stability

KR1: Demonstrating type hierarchies

In this section, we will explore the different subtypes of the abstract datatype `Number`. Starting from `Number`, we can see its subtypes through the function `subtypes()`.

```
In [2]: subtypes(Number)

Out[2]: 2-element Vector{Any}:
  Complex
  Real

Number is divided into two abstract subtypes: Complex for complex number types and Real for real number types. We can further check the subtypes in Real to identify the specific real number types.

In [3]: subtypes(Real)

Out[3]: 4-element Vector{Any}:
  AbstractFloat
  AbstractIrrational
  Integer
  Rational

In [4]: subtypes(Integer)

Out[4]: 3-element Vector{Any}:
  Bool
  Signed
  Unsigned

In [5]: subtypes(Signed)

Out[5]: 6-element Vector{Any}:
  BigInt
  Int128
  Int16
  Int32
  Int64
  Int8

In [6]: subtypes(BigInt)

Out[6]: Type[]

Thus if we trace the type hierarchy of BigInt, we can see that BigInt belongs to Signed integer types, which then belongs to the larger Integer abstract type, which then belongs to the Real number type.

Note as well that the Complex number type has no subtypes under it. (Complex is of the type UnionAll, which is similar to the user-defined struct.)

In [7]: subtypes(Complex)

Out[7]: Type[]

Using the function supertype(), we can trace back the Complex type to the Number abstract type.

In [8]: supertype(Complex)

Out[8]: Number
```

KR2: struct construction

For this section, we will use the `struct` construction of Julia to create composite types. `struct` instances usually are composed of at least two elements, each with its own fundamental type. Types assigned to each element can either be arbitrary or restricted, depending on how they are defined.

Here, we can construct a struct called `Planet`. This struct takes the following parameters: a string `planet_name` for the name of the planet, a number (which we will fix as a `Real` for our purposes) `planet_mass` for the mass of the planet, and a 2-D vector `planet_position` set relative to some solar position as the origin. The new type `Planet` is defined below.

```
In [9]: struct Planet
  planet_name::String
  planet_mass::Real #in kg
  planet_position::Vector #in AU
end
```

From here, we can instantiate multiple `Planet` objects.

```
In [10]: earth = Planet("Earth", 5.9722e24, [0.0, 1.0]);
mars = Planet("Mars", 6.39e23, [0.0, 1.5]);
mercury = Planet("Mercury", 3.285e23, [0.0, 0.4]);

println("The third planet from the Sun is $(earth.planet_name).")
println("$$(mars.planet_name) has a mass of $(mars.planet_mass) kg.")
println("$(mercury.planet_name) has an average distance of $(mercury.planet_position[2]) AU from the Sun.")

The third planet from the Sun is Earth.
Mars has a mass of 6.39e23 kg.
Mercury has an average distance of 0.4 AU from the Sun.
```

As demonstrated above, we can easily call elements from each `Planet` type object. We cannot, however, change the values of the elements, save for the vector `planet_position`.

```
In [11]: earth.planet_name = "Venus"

setfield! immutable struct of type Planet cannot be changed

Stacktrace:
 [1] setproperty!(x::Planet, f::Symbol, v::String)
    @ Base ./Base.jl:34
 [2] top-level scope
    @ In[11]:1
 [3] eval
    @ ./boot.jl:360 [inlined]
 [4] include_string(mapexpr::typeof(REPL.softscope), mod::Module, code::String, filename::String)
    @ Base ./loading.jl:1116

In [12]: earth.planet_position[:] = [1.0, 0.0]
```

```
Out[12]: 2-element Vector{Float64}:
 1.0
 0.0

In [13]: earth

Out[13]: Planet("Earth", 5.9722e24, [1.0, 0.0])
```

Note as well that `planet_position` is fixed as a 2-D vector, and its dimensions cannot be changed for this type of struct.

If we want a more flexible Type to use, we can opt to make a mutable struct. Let us define struct `CelestialObj` with the following elements.

```
In [14]: mutable struct CelestialObj
  obj_name::String
  obj_mass::Real
  obj_pos::Vector
end
```

Let us now define the object `mars_mutable` using this new struct.

```
In [15]: mars_mutable = CelestialObj("Mars", 6.39e23, [0.0, 1.5])

Out[15]: CelestialObj("Mars", 6.39e23, [0.0, 1.5])
```

Because `CelestialObj` is mutable, we can change the values of the elements freely (as long as it doesn't move out of their abstract supertype).

```
In [16]: mars_mutable.obj_name = "Mutated Mars";
mars_mutable.obj_mass = π;
mars_mutable.obj_pos = [16, 10, 7]; #taken from the current right ascension of Mars at J2000.0

mars_mutable

Out[16]: CelestialObj("Mutated Mars", π, [16, 10, 7])
```

It can be clearly seen that while mutability of the struct gives us some flexibility in terms of inputs, it can easily lead to issues such as type instability if not constrained properly.

We can also instead opt to parameterize some of the types in the struct to prevent overt type instability. A parameterized form of the non-mutable `Planet` struct would be

```
In [17]: struct PPlanet{T}
  planet_name::String
  planet_mass::T
  planet_pos::Vector{T}
end

In [18]: mercury_param = PPlanet{Float64}("Mercury", 3.285e23, [0, 0.4])

Out[18]: PPlanet{Float64}("Mercury", 3.285e23, [0.0, 0.4])
```

As we can see, specifying the type `T` through the parameterization fixes all elements taking the type `T` as `Float64`.

KR3: Demonstrating type inference with generator expressions

As demonstrated in the previous sessions, Julia has the natural feature of inferring the types of inputs it is given. Take for example the following generator expression

```
In [19]: [(x+1) for x in 1:5]

Out[19]: 5-element Vector{Int64}:
 2
 3
 4
 5
 6
```

which takes the integers `1` to `5` and adds by the integer `1`. Compare it with a similar generator expression which takes `floats` `1.0` to `5.0` and adds them to the `integer` `1`.

```
In [20]: [(x+1) for x in 1.0:5.0]

Out[20]: 5-element Vector{Float64}:
 2.0
 3.0
 4.0
 5.0
 6.0
```

We see for this particular generator that Julia immediately recognizes the input being generated as floats, and does the type promotion accordingly to allow the operation to be valid.

KR4-5: Type instabilities

Let `realroot(x)` be a function which accepts real number inputs `x` and outputs their square root when `x` is greater than zero, and outputs zero otherwise. This can be coded as

```
In [21]: realroot(x::Real) = x > 0 ? sqrt(x) : 0

Out[21]: realroot (generic function with 1 method)
```

The following examples demonstrate that the function does indeed produce the square roots of its positive inputs, and outputs zero for non-positive inputs.

```
In [22]: realroot(4)

Out[22]: 2.0

In [23]: realroot(4.0)

Out[23]: 2.0

In [24]: realroot(-4.0)

Out[24]: 0
```

We can see here however that there is already type instability inherent in this implementation: whereas `sqrt(x)` always outputs a `Float` regardless of whether `x` is inferred to be `Int` or `Float`, the other possible value `0` is strictly an `Int` regardless of the type of `x`. That is, both `Int` and `Float` inputs can result in either an `Int` or a `Float` output. The `@code_warntype` macro shows this problem clearly.

```
In [25]: @code_warntype realroot(4)

Variables
#self#::Core.Const{realroot}
x::Int64

Body::Union{Float64, Int64}
1 ┌ %1 = (x > 0)::Bool
  │   └─ goto #3 if not %1
2 ┌ %3 = Main.sqrt(x)::Float64
  │   └─ return %3
3 └─ return 0

In [26]: @code_warntype realroot(-4.0)

Variables
#self#::Core.Const{realroot}
x::Float64

Body::Union{Float64, Int64}
1 ┌ %1 = (x > 0)::Bool
  │   └─ goto #3 if not %1
2 ┌ %3 = Main.sqrt(x)::Float64
  │   └─ return %3
3 └─ return 0
```

Since `sqrt(x)` always outputs `Float` values in the valid domain, we can fix the type instability by forcing `0` to be a `Float` as well (set it to `0.0`).

```
In [27]: realroot_fixed(x::Real) = x > 0 ? sqrt(x) : 0.0

Out[27]: realroot_fixed (generic function with 1 method)
```

```
In [28]: @code_warntype realroot_fixed(4)

Variables
#self#::Core.Const{realroot_fixed}
x::Int64

Body::Float64
1 ┌ %1 = (x > 0)::Bool
  │   └─ goto #3 if not %1
2 ┌ %3 = Main.sqrt(x)::Float64
  │   └─ return %3
3 └─ return 0.0

In [29]: @code_warntype realroot_fixed(-4.0)

Variables
#self#::Core.Const{realroot_fixed}
x::Float64

Body::Float64
1 ┌ %1 = (x > 0)::Bool
  │   └─ goto #3 if not %1
2 ┌ %3 = Main.sqrt(x)::Float64
  │   └─ return %3
3 └─ return 0.0
```

As we can see, the type ambiguity and instability disappears after fixing all output types to `Float`.

KR6: Type ambiguity in Array operations

In general, one of the keys to optimizing code in Julia is to keep variable types as consistent and specific as possible so that the compiler does not spend time performing type inferencing at each operation. We can show the difference in runtimes with the following example.

Let `X` and `Y` be the following matrices:

```
In [30]: X = Float64[1 2 3
 4 5 6
 7 8 9];
Y = Real[1 2 3
 4 5 6
 7 8 9];
```

where the elements of `X` are all fixed as `Float64` values while the elements of `Y` are arbitrarily typed as `Real`. We can verify that these matrices are still of the type `Array` by checking their type.

```
In [31]: typeof(X)

Out[31]: Matrix{Float64} (alias for Array{Float64, 2})

In [32]: typeof(Y)

Out[32]: Matrix{Real} (alias for Array{Real, 2})
```

We can take the determinant of both matrices

```
In [33]: det(X)

Out[33]: 0.0

In [34]: det(Y)

Out[34]: 6.661338147750939e-16

In [35]: typeof(det(X))

Out[35]: Float64

In [36]: typeof(det(Y))

Out[36]: Float64
```

and we see that the determinant of `X` is zero, while the determinant of `Y` is very small but nonzero. If done by hand, the determinant of the above matrix is exactly zero, which seems to suggest that the type ambiguity allowed by in `Y` has resulted in some rounding errors during the `det()` operation. No such errors are present for the purely `Float64` matrix.

We can also compare the benchmarked times for both `X` and `Y` inputs.

```
In [37]: xdetbench = @benchmark det($X)

Out[37]: BenchmarkTools.Trial: 10000 samples with 308 evaluations.
 Range (min ... max): 266.211 ns ... 4.935 μs      GC (min ... max): 0.00% ... 91.10%
 Time (median): 319.005 ns      GC (median): 0.00%
 Time (mean ± σ): 334.254 ns ± 161.406 ns      GC (mean ± σ): 1.98% ± 3.88%

Histogram: log(frequency) by time
266 ns                                     513 ns <

Memory estimate: 272 bytes, allocs estimate: 2.
```

```
In [38]: ydetbench = @benchmark det($Y)

Out[38]: BenchmarkTools.Trial: 10000 samples with 9 evaluations.
 Range (min ... max): 2.253 μs ... 6.103 μs      GC (min ... max): 0.00% ... 0.00%
 Time (median): 2.544 μs      GC (median): 0.00%
 Time (mean ± σ): 2.629 μs ± 391.576 ns      GC (mean ± σ): 0.00% ± 0.00%

Histogram: frequency by time
2.25 μs                                     4.59 μs <

Memory estimate: 608 bytes, allocs estimate: 23.
```

```
In [39]: medianratio = median(Ydetbench.times)/median(Xdetbench.times);
println("det(X::Float64) is $(round(medianratio; digits = 2)) times faster than det(Y)

det(X::Float64) is 7.97 times faster than det(Y::Real).
```

As we can see, not only does clearer typing provide better numerical accuracy, but it also improves code runtime.